

Were We Already There? Applying Minimal Generalization to the SIGMORPHON-UniMorph Shared Task on Cognitively Plausible Morphological Inflection

Colin Wilson

Johns Hopkins University
colin.wilson@jhu.edu

Second Author

Affiliation / Address line 1
Affiliation / Address line 2
Affiliation / Address line 3
email@domain

Abstract

xxx

1 Introduction

In a landmark paper, Albright and Hayes (2003) proposed a model that learns morphological rules by recursive **minimal generalization** from lexeme-specific examples (e.g., $\text{ɪ} \rightarrow \text{ʌ} / \text{st} _ \eta$ for *sting* \sim *stung* and $\text{ɪ} \rightarrow \text{ʌ} / \text{fl} _ \eta$ for *fling* \sim *flung* generalized to $\text{ɪ} \rightarrow \text{ʌ} / \text{X} [-\text{syllabic}, +\text{coronal}, +\text{anterior}, \dots] _ \eta$).¹ The model was presented more formally in Albright and Hayes (2002), along with evidence that the rules it learns for the English past tense give a good account of native speakers' productions and ratings in wug-test experiments (e.g. judgments that *splung* is quite acceptable as the past tense of the novel verb *spling*). In addition to providing further analysis of the experimental data, Albright and Hayes (2003) compared their proposal with early connectionist models of morphology (e.g., Plunkett and Juola, 1999) and an analogical or 'family resemblance' model inspired by research on psychological categories (Nakisa et al., 2001).

Along with Albright (2002), which presents a parallel treatment of Italian inflection, Albright & Hayes's study of the English past tense is a paradigm example of theory-driven, multiple-methodology, open and reproducible research in cognitive science.² Their model has enduring significance for the study of morphological learning

and productivity in English (e.g., Rácz et al., 2014, 2020; Corkery et al., 2019) and many other languages (e.g., Hijazi Arabic: Ahyad 2019; Japanese: Oseki et al. 2019; Korean: Albright and Kang 2009; Navajo: Albright and Hayes 2006; Portuguese: Veríssimo and Clahsen 2014; Russian: Kapatsinski 2010; Tgdaya Seediq: Kuo 2020; Spanish: Albright and Hayes 2003; Swedish: Strik 2014).

In this paper, we apply a partial reimplementa-tion of the Albright & Hayes model to wug-test rating data from three languages (Dutch, English, and German) collected for the SIGMORPHON-UniMorph 2021 Shared Task. Our version of the model is based purely on minimal generalization of morphological rules, as described in §3.1 of Albright and Hayes (2002) and reviewed below. It does not include additional mechanisms for learning phonological rules, and expanding or reigning in morphological rules, that were part of the original model (see Albright and Hayes, 2002, §3.3 - §3.7). We think there is much to be gained from considering minimal generalization on its own, with the other mechanisms ablated away, as borne out by competitive results on the shared task. For convenience, we refer to the original model as **min-gen** and our bare-bones version as **mingen0**.

1.1 Outline

2 Model Description

2.1 Inputs

The model takes as input a set of wordform pairs, one per lexeme, that instantiate the same morphological relationship. In simulations of English past tense formation, these are pairs of bare verb stems and past tense forms such as $\langle \text{ʌwɔk}, \text{ʌwɔkt} \rangle$, $\langle \text{tɔk}, \text{tɔkt} \rangle$, $\langle \text{stɪŋ}, \text{stɪŋ} \rangle$, $\langle \text{flɪŋ}, \text{flɪŋ} \rangle$, and $\langle \text{kæt}, \text{kæt} \rangle$. Wordforms consist of phonological segments (here, in broad transcription) delimited by special beginning and end of string symbols. The set Σ of phonological seg-

¹The square brackets contain the shared phonological feature specifications of /t/ and /l/, which in the feature system used here are xxx.

²Albright & Hayes released both the results of their wug-test experiments and an implementation of their model (visit <http://www.mit.edu/~albright/mgl/> and <https://linguistics.ucla.edu/people/hayes/RulesVsAnalogy/index.html>). An impediment to large-scale simulation with the model is that it runs from a GUI interface only. As part of the present project, we have added a command line interface to the original source code (available on request).

ments, and the set $\Sigma_{\#} = \Sigma \cup \{\bowtie, \bowtie\}$, are assumed to be given to the model.

The model also requires a phonological feature specification for each of the symbols that appears in wordforms. We used a well-known feature set, augmenting it with orthogonal feature specifications for the delimiters \bowtie and \bowtie .³ Φ is the set of all possible (partial) feature specifications over the chosen set and $\phi(x)$ gives the features of $x \in \Sigma_{\#}$.

2.2 Base rules

For each wordform pair, the model constructs a lexeme-specific morphological rule by first identifying the longest common prefix excluding \bowtie (C), then the longest common suffix from the remainder (D), and finally identifying the remaining symbols in the first (A) and second (B) members of the pair. The symbol $\emptyset \notin \Sigma_{\#}$ denotes the empty string.⁴ The rule formed from $\langle \bowtie w \bowtie k \bowtie, \bowtie w \bowtie k t \bowtie \rangle$ has the components $C = \bowtie w \bowtie k$, $D = \bowtie$, $A = \emptyset$ and $B = t$ (i.e., $\emptyset \rightarrow t / \bowtie w \bowtie k _ \bowtie$). The rule for $\langle \bowtie k \bowtie t \bowtie, \bowtie k \bowtie t \bowtie \rangle$ is $\emptyset \rightarrow \emptyset \bowtie k \bowtie t _ \bowtie$.

2.3 Minimal Generalization

Given any two base rules R_1 and R_2 that make the same change ($A \rightarrow B$), the model forms a possibly more general rule by aligning and comparing their contexts. The minimal generalization operation, $R = R_1 \sqcap R_2$, carries over the common change of the two base rules and applies independently to their left-hand (C_1, C_2) and right-hand (D_1, D_2) contexts. For convenience, we define minimal generalization of the right-hand contexts. Minimal generalization of the left-hand contexts can be performed by reversing C_1 and C_2 , applying the definition for right-hand contexts, and reversing the result.

The minimal generalization $D = D_1 \sqcap D_2$ is defined procedurally by first extracting the longest common prefix ($\sigma_{1 \wedge 2}$) of the two contexts and then operating on the remainders (D'_1, D'_2). If both D'_1 and D'_2 are empty then $D = \sigma_{1 \wedge 2}$. If one but not both of them are empty then $D = \sigma_{1 \wedge 2} X$, where $X \notin \Sigma_{\#}$ is a variable over symbol sequences (i.e.,

X stands for $\Sigma_{\#}^*$). If neither is empty, then the operation determines whether their initial symbols have any shared features; for this purpose it is useful to consider $\phi(x)$ as a function from symbols to sets of feature-value pairs, in which case the common features are found by set intersection.

If there are no common features, $\phi_{1 \wedge 2} = \emptyset$, then as before $D = \sigma_{1 \wedge 2} X$. Otherwise, the set of common features $\phi_{1 \wedge 2} \neq \emptyset$ is appended to $\sigma_{1 \wedge 2}$, the first symbol is removed from D'_1 and D'_2 , and the operation applies to the remainders. If both remainders are empty then $D = \sigma_{1 \wedge 2} \phi_{1 \wedge 2}$, otherwise $D = \sigma_{1 \wedge 2} \phi_{1 \wedge 2} X$.

In summary, the generalized right-hand context D consists of the longest common prefix shared by D_1 and D_2 , followed by a single set of shared features (if any), followed by X in case there are no shared features or one context is longer than the other. With the change and generalized left-hand context C determined as already described, the result of applying minimal generalization to the two base rules is $R = A \rightarrow B/C _ D$.⁵

2.4 Recursive Minimal Generalization

Let \mathcal{R}_1 be the set of base rules (one per wordform pair in the input data) and \mathcal{R}_2 be the set containing all of the base rules and the result of applying minimal generalization to each eligible pair of base rules. While the rules of \mathcal{R}_2 have greater collective scope, they are nevertheless unlikely to account for the level of morphological productivity shown by native speakers. For example, English speakers can systematically rate and produce past tense forms of novel verbs that contain unusual segment sequences, such as *ploomf* /*ploomf*/ (e.g., [Prasada and Pinker, 1993](#)). Albright & Hayes propose to apply minimal generalization recursively and demonstrate that this can yield rules of great generality (e.g., in our notation, $\emptyset \rightarrow t / X [-\text{voice}] _ \bowtie$).

In the original proposal, recursive minimal generalization was defined only for pairs that include one base rule; it was conjectured that no additional generalizations could result from dropping this restriction. Here we define the operation for any

³The phonological feature specifications are available from Bruce Hayes's website, <https://linguistics.ucla.edu/people/hayes/120a/Index.htm#features>. xxx binary with 0s xxx original features distributed with the model included scalar features xxx for example. Alternative binary (with underspecification) feature sets are xxx phoible xxx panphon.

⁴In other common notations, the empty string is denoted by λ . xxx by construction can occur only in A or B

⁵There could be a small difference between our definition of context generalization and that in [Albright and Hayes \(2002\)](#), hinging on whether the empty feature set is allowed in rules. In our definition, $\phi_{1 \wedge 2} = \emptyset$ is replaced by the variable X . It is possible that the original proposal intended for empty and non-empty feature sets to be treated alike. The definitions can diverge when applied to contexts that are of identical length and share all but the last (resp. first) segments, in which case our version would result in a broader rule.

two right-hand contexts $D_1, D_2 \in \Sigma_{\#}^*(\Phi)(X)$. As before, only rules that make the same change are eligible for generalization and the operation applies to left-hand contexts under reversal.

The revised definition of $D = D_1 \sqcap D_2$ is identical to that given above except that we must consider input contexts that contain feature sets and X (which previously could occur only in outputs). As before, we first identify the longest common prefix of symbols from $\Sigma_{\#}$ in the two contexts, $\sigma_{1 \wedge 2}$, and then operate on the remainders (D'_1, D'_2). If both D'_1 and D'_2 are empty then $D = \sigma_{1 \wedge 2}$. If one but not both of them are empty then $D = \sigma_{1 \wedge 2}X$. If both are non-empty then their initial elements are either symbols in $\Sigma_{\#}$, feature sets in Φ , or X . Replace any initial symbol $x \in \Sigma_{\#}$ with its feature set $\phi(x)$, extend the function ϕ so that $\phi(X) = \emptyset$, and compute the unification $\phi_{1 \cap 2}$ of the initial elements. The rest of the definition is unchanged (see the end of §2.3).

By construction, the contexts that result from this operation are also in $\Sigma_{\#}^*(\Phi)(X)$ (*i.e.*, no segment can occur after a feature set, there is at most one feature set, X can occur only at the end of the context, etc.). Therefore, the revised definition supports the application of minimal generalization to its own products. Let \mathcal{R}_k be the set of rules containing every member of \mathcal{R}_{k-1} and the result of applying minimal generalization to each eligible pair of rules in \mathcal{R}_{k-1} (for $k > 1$). In principle, there is an infinite sequence of rules sets related by inclusion $\mathcal{R}_1 \subseteq \mathcal{R}_2 \subseteq \mathcal{R}_3 \dots$. In practice, the equality becomes strict after a small number of iterations of minimal generalization (typically 6-7), at which point there are no more rules to be found.

2.5 Completeness

Having defined minimal generalization for arbitrary contexts (as allowed by the model), we can revisit the conjecture that nothing is lost by restricting the operation to pairs at least one of which is a base rule. This is a practical concern, as the number of base rules is a constant determined by the input data while the number of generalized rules (for a given change) could increase exponentially. Intuitively, each rule learned by the model, no matter how general, has a (non-unique) ‘history’ of base rules from which it was derived.

2.6 Scoring

xxx

2.7 Pruning

xxx

3 Results

4 Conclusions and Future Directions

These instructions are for authors submitting papers to *ACL conferences using L^AT_EX. They are not self-contained. All authors must follow the general instructions for *ACL proceedings,⁶ and this document contains additional instructions for the L^AT_EX style files.

The templates include the L^AT_EX source of this document (`acl.tex`), the L^AT_EX style file used to format it (`acl.sty`), an ACL bibliography style (`acl_natbib.bst`), an example bibliography (`custom.bib`), and the bibliography for the ACL Anthology (`anthology.bib`).

5 Engines

To produce a PDF file, pdfL^AT_EX is strongly recommended (over original L^AT_EX plus dvips+ps2pdf or dvi2pdf). XeL^AT_EX also produces PDF files, and is especially suitable for text in non-Latin scripts.

6 Preamble

The first line of the file must be

```
\documentclass[11pt]{article}
```

To load the style file in the review version:

```
\usepackage[review]{acl}
```

For the final version, omit the `review` option:

```
\usepackage{acl}
```

To use Times Roman, put the following in the preamble:

```
\usepackage{times}
```

(Alternatives like `txfonts` or `newtx` are also acceptable.)

Please see the L^AT_EX source of this document for comments on other packages that may be useful.

Set the title and author using `\title` and `\author`. Within the author list, format multiple authors using `\and` and `\And` and `\AND`; please see the L^AT_EX source for examples.

By default, the box containing the title and author names is set to the minimum of 5 cm. If you need more space, include the following in the preamble:

```
\setlength\titlebox{<dim>}
```

where `<dim>` is replaced with a length. Do not set this length smaller than 5 cm.

⁶<http://acl-org.github.io/ACL/PUB/formatting.html>

Command	Output	Command	Output
<code>\a</code>	ä	<code>\c c</code>	ç
<code>\^e</code>	ê	<code>\u g</code>	ğ
<code>\i</code>	ì	<code>\l</code>	ł
<code>\.I</code>	İ	<code>\~n</code>	ñ
<code>\o</code>	ø	<code>\H o</code>	ő
<code>\'u</code>	ú	<code>\v r</code>	ř
<code>\aa</code>	å	<code>\ss</code>	ß

Table 1: Example commands for accented characters, to be used in, e.g., BibT_EX entries.

7 Document Body

7.1 Footnotes

Footnotes are inserted with the `\footnote` command.⁷

7.2 Tables and figures

See Table 1 for an example of a table and its caption. **Do not override the default caption sizes.**

7.3 Hyperlinks

Users of older versions of L^AT_EX may encounter the following error during compilation:

```
\pdfendlink ended up in
different nesting level
than \pdfstartlink.
```

This happens when pdfL^AT_EX is used and a citation splits across a page boundary. The best way to fix this is to upgrade L^AT_EX to 2018-12-01 or later.

7.4 Citations

Table 2 shows the syntax supported by the style files. We encourage you to use the `natbib` styles. You can use the command `\citet` (cite in text) to get “author (year)” citations, like this citation to a paper by ?. You can use the command `\citep` (cite in parentheses) to get “(author, year)” citations (?). You can use the command `\citealp` (alternative cite without parentheses) to get “author, year” citations, which is useful for using citations within parentheses (e.g. ?).

7.5 References

The L^AT_EX and BibT_EX style files provided roughly follow the American Psychological Association format. If your own bib file is named `custom.bib`, then placing the following before any appendices in

⁷This is a footnote.

Output	natbib command	Old ACL-style command
(?)	<code>\citep</code>	<code>\cite</code>
?	<code>\citealp</code>	no equivalent
?	<code>\citet</code>	<code>\newcite</code>
(?)	<code>\citeyearpar</code>	<code>\shortcite</code>

Table 2: Citation commands supported by the style file. The style is based on the natbib package and supports all natbib citation commands. It also supports commands defined in previous ACL style files for compatibility.

your L^AT_EX file will generate the references section for you:

```
\bibliographystyle{acl_natbib}
\bibliography{custom}
```

You can obtain the complete ACL Anthology as a BibT_EX file from <https://aclweb.org/anthology/anthology.bib.gz>. To include both the Anthology and your own .bib file, use the following instead of the above.

```
\bibliographystyle{acl_natbib}
\bibliography{anthology, custom}
```

Please see Section 8 for information on preparing BibT_EX files.

7.6 Appendices

Use `\appendix` before any appendix section to switch the section numbering over to letters. See Appendix A for an example.

8 BibT_EX Files

Unicode cannot be used in BibT_EX entries, and some ways of typing special characters can disrupt BibT_EX’s alphabetization. The recommended way of typing special characters is shown in Table 1.

Please ensure that BibT_EX records contain DOIs or URLs when possible, and for all the ACL materials that you reference. Use the `doi` field for DOIs and the `url` field for URLs. If a BibT_EX entry has a URL or DOI field, the paper title in the references section will appear as a hyperlink to the paper, using the `hyperref` L^AT_EX package.

Acknowledgements

This research was partially supported by NSF grant BCS-1844780.

References

- Honaida Yousuf Ahyad. 2019. *Vowel Distribution in the Hijazi Arabic Root*. Ph.D., State University of New York at Stony Brook, United States – New York.
- Adam Albright. 2002. Islands of Reliability for Regular Morphology: Evidence from Italian. *Language*, 78(4):684–709.
- Adam Albright and Bruce Hayes. 2002. [Modeling english past tense intuitions with minimal generalization](#). In *Proceedings of the ACL-02 Workshop on Morphological and Phonological Learning -*, volume 6, pages 58–69, Not Known. Association for Computational Linguistics.
- Adam Albright and Bruce Hayes. 2003. [Rules vs. analogy in English past tenses: A computational/experimental study](#). *Cognition*, 90(2):119–161.
- Adam Albright and Bruce Hayes. 2006. Modeling productivity with the Gradual Learning Algorithm: The problem of accidentally exceptionless generalizations. In Gisbert Fanselow, Caroline Fery, Matthias Schlesewsky, and Ralf Vogel, editors, *Gradience in Grammar: Generative Perspectives*, pages 185–204. Oxford University Press, Oxford.
- Adam Albright and Yoonjung Kang. 2009. Predicting innovative alternations in Korean verb paradigms. *Current issues in unity and diversity of languages: Collection of the papers selected from the CIL 18, held at Korea University in Seoul*, pages 1–20.
- Benjamin Börschinger and Mark Johnson. 2011. [A particle filter algorithm for Bayesian wordsegmentation](#). In *Proceedings of the Australasian Language Technology Association Workshop 2011*, pages 10–18, Canberra, Australia.
- Maria Corkery, Yevgen Matushevych, and Sharon Goldwater. 2019. [Are we there yet? encoder-decoder neural networks as cognitive models of English past tense inflection](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3868–3877, Florence, Italy. Association for Computational Linguistics.
- James Goodman, Andreas Vlachos, and Jason Naradowsky. 2016. [Noise reduction and targeted exploration in imitation learning for Abstract Meaning](#)

[Representation parsing](#). In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1–11, Berlin, Germany. Association for Computational Linguistics.

Mary Harper. 2014. [Learning from 26 languages: Program management and science in the babel program](#). In *Proceedings of COLING 2014, the 25th International Conference on Computational Linguistics: Technical Papers*, page 1, Dublin, Ireland. Dublin City University and Association for Computational Linguistics.

Vsevolod Kapatsinski. 2010. [Velar palatalization in Russian and artificial grammar: Constraints on models of morphophonology](#). *Laboratory Phonology*, 1(2):361–393.

Jennifer Kuo. 2020. *Evidence for Base-Driven Alternation in Tgdaya Seediq*. Ph.D. thesis, UCLA.

R. C. Nakisa, Kim Plunkett, and U Hahn. 2001. A cross-linguistic comparison of single and dual-route models of inflectional morphology. In P Broder and J Murre, editors, *Models of Language Acquisition: Inductive and Deductive Approaches*. MIT Press, Cambridge, MA.

Yohei Oseki, Yasutada Sudo, Hiromu Sakai, and Alec Marantz. 2019. [Inverting and modeling morphological inflection](#). In *Proceedings of the 16th Workshop on Computational Research in Phonetics, Phonology, and Morphology*, pages 170–177, Florence, Italy. Association for Computational Linguistics.

Kim Plunkett and Patrick Juola. 1999. [A Connectionist Model of English Past Tense and Plural Morphology](#). *Cognitive Science*, 23(4):463–490.

S. Prasada and Steven Pinker. 1993. Generalisation of regular and irregular morphological patterns. *Language and cognitive processes*, 8(1):1–56.

Péter Rácz, Clay Beckner, Jennifer B. Hay, and Janet B. Pierrehumbert. 2020. [Morphological convergence as on-line lexical analogy](#). *Language*, 96(4):735–770.

Péter Rácz, Clayton Beckner, Jennifer B. Hay, and Janet B. Pierrehumbert. 2014. [Rules, analogy, and social factors codetermine past-tense formation patterns in English](#). In *Proceedings of the 2014 Joint Meeting of SIGMORPHON and SIGFSM*, pages 55–63, Baltimore, Maryland. Association for Computational Linguistics.

Oscar Strik. 2014. [Explaining tense marking changes in Swedish verbs: An application of two analogical computer models](#). *Journal of Historical Linguistics*, 4(2):192–231.

João Veríssimo and Harald Clahsen. 2014. [Variables and similarity in linguistic generalization: Evidence from inflectional classes in Portuguese](#). *Journal of Memory and Language*, 76:61–79.

A Example Appendix

This is an appendix.